

# Introduction

This document briefly describes how to access the Hadoop cluster setup for COMP 421 and submit jobs using pig and associated basic monitoring. This is a growing document that will be revised as needed to elaborate the information provided as necessity arises, hence we recommend you to cross check the version number of this document with the one in mycourses which will always be the most up to date.

Please read this document completely before you start submitting your pig scripts !

If you copy-paste any instructions from this document, double check that the pasted text matches what is there in this document, letters like hyphen(-) , quotation (') marks etc usually get translated wrong.

## Environment Setup

You will have to login to `winter2022-comp421.cs.mcgill.ca` using your SOCS account to write pig scripts. Additionally, you need to include `/data/cs421/softwares/apache/pig-0.15.0/bin` in your PATH.

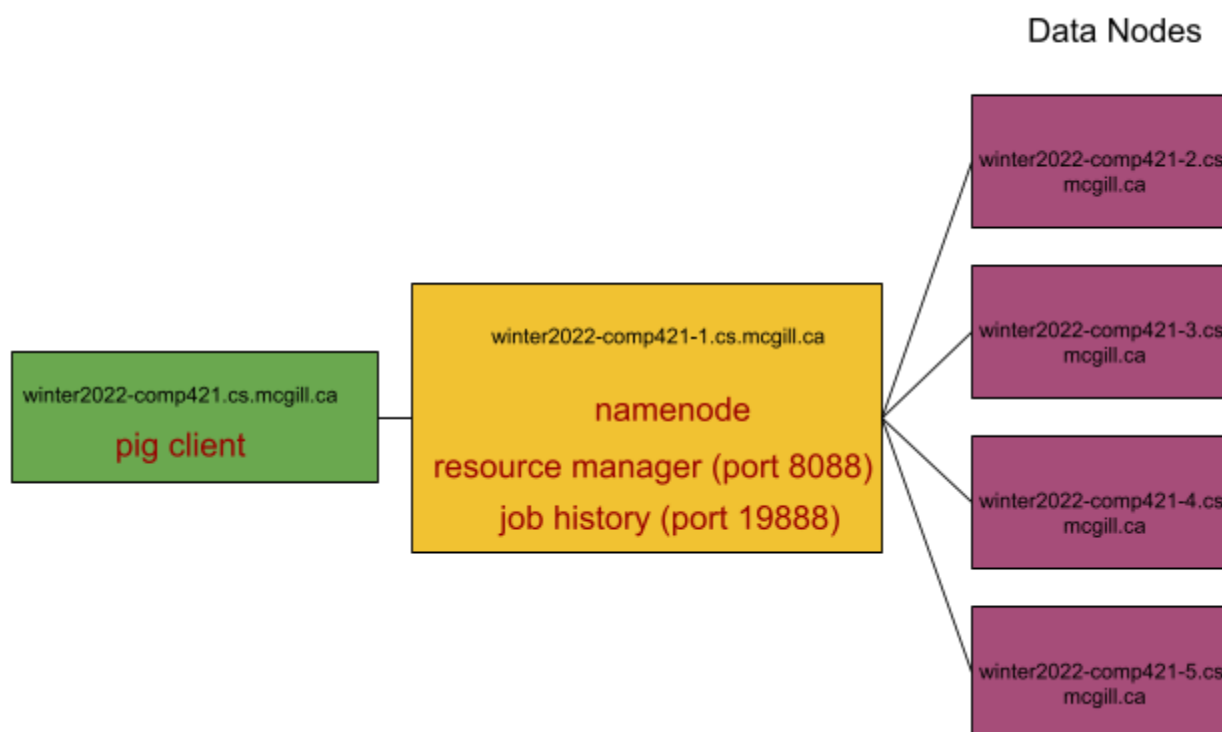
```
$ PATH=/data/cs421/softwares/apache/pig-0.15.0/bin:$PATH
```

so that when you type

```
$ which pig
```

```
/data/cs421/softwares/apache/pig-0.15.0/bin/pig
```

You get the path to the pig executable.



The hadoop cluster consists of one name node (winter2022-comp421-1) and four data nodes (winter2022-comp421-2 ... winter2022-comp421-5)

Please refrain from attempting to login on to the cluster nodes directly or writing scripts there. The user home filesystems in these nodes are temporary and the files created in them will disappear on system reboots !! you also do not have access to run pig in these nodes.

## Execution and Monitoring

You can execute the pig commands by either writing them into a script and passing it as an argument to the pig commands

```
pig example.pig
```

Or by just typing pig, getting the grunt prompt and then typing in each command on the grunt prompt.

```
$ pig
```

```
22/03/03 17:11:06 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
22/03/03 17:11:06 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
22/03/03 17:11:06 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
```

```
2022-03-03 17:11:06,612 [main] INFO    org.apache.pig.Main - Apache Pig version 0.15.0 (r1682971) compiled Jun
01 2015, 11:44:35
2022-03-03 17:11:06,612 [main] INFO    org.apache.pig.Main - Logging error messages to:
/home/2013/jdsilv2/MyStuff/hd-2022/pig_1614809466610.log
2022-03-03 17:11:06,636 [main] INFO    org.apache.pig.impl.util.Utils - Default bootup file
/home/2013/jdsilv2/.pigbootup not found
2022-03-03 17:11:07,139 [main] INFO    org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is
deprecated. Instead, use mapreduce.jobtracker.address
2022-03-03 17:11:07,139 [main] INFO    org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is
deprecated. Instead, use fs.defaultFS
2022-03-03 17:11:07,139 [main] INFO    org.apache.pig.backend.hadoop.executionengine.HExecutionEngine -
Connecting to hadoop file system at: hdfs://winter2022-comp421-1.cs.mcgill.ca:9000
2022-03-03 17:11:07,947 [main] INFO    org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is
deprecated. Instead, use fs.defaultFS
grunt>
```

As you can see pig will output a lot of information whether you execute commands from grunt or through a pig script. Additionally, it will also capture error messages and output them into a log file which is useful for debugging later (in the above example this has been highlighted in red) . **These messages produced in the pig output will be your primary source for information and debugging** as most of the errors will be intercepted by the pig and will have to do with errors in pig syntax / semantics Hence you may not find any info about these on the job history logs of Hadoop as the pig would not even have submitted the job yet.

Among many informative output by pig is the SimplePigStats that tell you how start and end time of the pig script, the number of jobs involved in it and the number of Maps/Reduces run on each job.

```
2016-03-30 15:53:02,417 [main] INFO    org.apache.pig.tools.pigstats.mapreduce.SimplePigStats - Script
Statistics:


HadoopVersion      PigVersion  UserIdStartedAt  FinishedAt  Features
2.7.2 0.15.0jdsilv2      2016-03-30 15:51:43      2016-03-30 15:53:02      ORDER_BY,FILTER,LIMIT

Success!

Job Stats (time in seconds):
JobId Maps   Reduces      MaxMapTime  MinMapTime  AvgMapTime  MedianMapTime      MaxReduceTime
MinReduceTime      AvgReduceTime      MedianReducetime  Alias Feature      Outputs
job_1458239741221_0063  1      0      4      4      4      4      0      0      0      0      fltrd,gen,raw
MAP_ONLY
job_1458239741221_0064  1      1      2      2      2      2      3      3      3      3      odred SAMPLER
job_1458239741221_0065  1      1      2      2      2      2      3      3      3      3      odred ORDER_BY,COMBINER
job_1458239741221_0066  1      1      3      3      3      3      3      3      3      3      odred
hdfs://winter2022-comp421-1.cs.mcgill.ca:9000/tmp/temp-1164992124/tmp1925912064,
```

While pig will output very detailed info into the terminal about the operations it is executing, you can also check the resource manager UI to see if your job is running or is in the pending queue (which can happen if there are too many jobs in the system).  
**ALL WEB ACCESS BELOW NEED TO GO THROUGH MCGILL VPN !!**

<http://winter2022-comp421-1.cs.mcgill.ca:8088/cluster/scheduler>



## NEW,NEW\_SAVING,SUBMITTED,ACCEPTED,RUNNING Applications

Cluster

About

Nodes

Node Labels

Applications

NEW

NEW\_SAVING

SUBMITTED

ACCEPTED

RUNNING

FINISHED

FAILED

KILLED

Scheduler

Tools

Cluster Metrics

Apps Submitted	Apps Pending	Apps Running	Apps Completed	Containers Running	Memory Used	Memory Total	Memory Reserved	VCores Used	VCores Total	VCores Reserved	Active Nodes	Decommissioned Nodes
18	0	2	16	2	4 GB	24 GB	0 B	2	24	0	3	0

Scheduler Metrics

Scheduler Type	Scheduling Resource Type	Minimum Allocation	Maximum Allocation
Capacity Scheduler	[MEMORY]	<memory:1024, vCores:1>	<memory:5120, vCores:1>

Application Queues

Legend:

Capacity

Used

Used (over capacity)

Max Capacity

root

default

system


comp421

Show 20 entries

ID	User	Name	Application Type	Queue	StartTime	FinishTime	State	FinalStatus	Progress
application_1458239741221_0018	jdsilv2	PigLatin:example.pig	MAPREDUCE	comp421	Tue Mar 29 13:43:09 -0400 2016	N/A	ACCEPTED	UNDEFINED	

Showing 1 to 1 of 1 entries

Further, you can use the job history UI to look at a more finer level of log messages (this is also the place you will have to go once the job is completed to check for any messages as the resource manager shows mostly information regarding jobs that are currently active). This page also displays how many maps / reduces were used for each of the jobs and is a good indication of parallelism. <http://winter2022-comp421-1.cs.mcgill.ca:19888/jobhistory>



JobHistory

▼ Application

About  
Jobs

► Tools

Retired Jobs

Show 20 entries


Search:

Submit Time	Start Time	Finish Time	Job ID	Name	User	Queue	State	Maps Total	Maps Completed	Reduces Total	Reduces Completed
2016.03.29 13:42:53 EDT	2016.03.29 13:43:01 EDT	2016.03.29 13:43:07 EDT	<a href="#">job_1458239741221_0017</a>	PigLatin:example.pig	jdsilv2	comp421	SUCCEEDED	1	1	0	0
2016.03.29 13:43:09 EDT	2016.03.29 13:43:17 EDT	2016.03.29 13:43:31 EDT	<a href="#">job_1458239741221_0018</a>	PigLatin:example.pig	jdsilv2	comp421	SUCCEEDED	1	1	1	1
2016.03.29 13:43:35 EDT	2016.03.29 13:43:40 EDT	2016.03.29 13:43:52 EDT	<a href="#">job_1458239741221_0019</a>	PigLatin:example.pig	jdsilv2	comp421	SUCCEEDED	1	1	1	1
2016.03.29 13:43:56 EDT	2016.03.29 13:44:02 EDT	2016.03.29 13:44:13 EDT	<a href="#">job_1458239741221_0020</a>	PigLatin:example.pig	jdsilv2	comp421	SUCCEEDED	1	1	1	1

Showing 1 to 4 of 4 entries

First Previous 1 Next Last

It is important to note that one pig script can result in multiple mapreduce jobs (for various steps in the pig script). You will find a Job ID for each one of them in the job history UI. You can click on the link for one of your jobs to receive additional information. (shown below)



MapReduce Job job\_1458239741221\_0017

► Application

▼ Job

Overview  
Counters  
Configuration  
Map tasks  
Reduce tasks

► Tools

Job Overview

Job Name: PigLatin:example.pig

User Name: jdsilv2

Queue: comp421

State: SUCCEEDED

Uberized: false

Submitted: Tue Mar 29 13:42:53 EDT 2016

Started: Tue Mar 29 13:43:01 EDT 2016

Finished: Tue Mar 29 13:43:07 EDT 2016

Elapsed: 6sec


Diagnostics:

Average Map Time 3sec

ApplicationMaster			
Attempt Number	Start Time	Node	Logs
1	Tue Mar 29 13:42:58 EDT 2016	cs421-hd5:8042	<a href="#">logs</a>

Task Type	Total	Complete	
Map	1	1	
Reduce	0	0	
Attempt Type	Failed	Killed	Successful
Maps	0	0	1
Reduces	0	0	0

You can click on the logs link for more log messages.



▼ Application

About  
Jobs

► Tools

Log Type: stderr

Log Upload Time: Tue Mar 29 13:43:14 -0400 2016

Log Length: 1703

Mar 29, 2016 1:43:00 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register

INFO: Registering org.apache.hadoop.mapreduce.v2.app.webapp.JAXBContextResolver as a provider class

Mar 29, 2016 1:43:00 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register

INFO: Registering org.apache.hadoop.yarn.webapp.GenericExceptionHandler as a provider class

Mar 29, 2016 1:43:00 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory register

INFO: Registering org.apache.hadoop.mapreduce.v2.app.webapp.AMWebServices as a root resource class

Mar 29, 2016 1:43:00 PM com.sun.jersey.server.impl.application.WebApplicationImpl \_initiate

INFO: Initiating Jersey application, version 'Jersey: 1.9 09/02/2011 11:17 AM'

Mar 29, 2016 1:43:00 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider

INFO: Binding org.apache.hadoop.mapreduce.v2.app.webapp.JAXBContextResolver to GuiceManagedComponentProvider with the scope "Singleton"

Mar 29, 2016 1:43:01 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider

INFO: Binding org.apache.hadoop.yarn.webapp.GenericExceptionHandler to GuiceManagedComponentProvider with the scope "Singleton"

Mar 29, 2016 1:43:01 PM com.sun.jersey.guice.spi.container.GuiceComponentProviderFactory getComponentProvider

INFO: Binding org.apache.hadoop.mapreduce.v2.app.webapp.AMWebServices to GuiceManagedComponentProvider with the scope "PerRequest"

log4j:WARN No appenders could be found for logger (org.apache.hadoop.ipc.Server).

log4j:WARN Please initialize the log4j system properly.

log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.

Log Type: stdout

Log Upload Time: Tue Mar 29 13:43:14 -0400 2016

Log Length: 0

Log Type: syslog

Log Upload Time: Tue Mar 29 13:43:14 -0400 2016

Log Length: 26260

Showing 4096 bytes of 26260 total. Click [here](#) for the full log.

ScheduledMaps:0 ScheduledReds:0 AssignedMaps:1 AssignedReds:0 CompletedMaps:1 CompletedReds:0 ContAlloc:1 ContRel:0 HostLocal:1 RackLocal:0

2016-03-29 13:43:08,027 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Copied to done location: hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

2016-03-29 13:43:08,030 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Copying hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

2016-03-29 13:43:08,069 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Copied to done location: hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

2016-03-29 13:43:08,072 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Moved tmp to done: hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

2016-03-29 13:43:08,074 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Moved tmp to done: hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

2016-03-29 13:43:08,075 INFO [eventHandlingThread] org.apache.hadoop.mapreduce.jobhistory.JobHistoryEventHandler: Moved tmp to done: hdfs://cs421-hd1.cs.mcgill.ca:9000/tmp/

## Things to know

HDFS will refuse to overwrite files, this can create issues if in your pig script you are using `STORE` commands. In order to delete any such files, one possible way is to start pig interactively and then use the `rm` command as illustrated below.

```
pig
16/03/29 14:58:54 INFO pig.ExecTypeProvider: Trying ExecType : LOCAL
16/03/29 14:58:54 INFO pig.ExecTypeProvider: Trying ExecType : MAPREDUCE
16/03/29 14:58:54 INFO pig.ExecTypeProvider: Picked MAPREDUCE as the ExecType
2016-03-29 14:58:54,161 [main] INFO org.apache.pig.Main - Apache Pig version 0.15.0 (r1682971) compiled Jun
01 2015, 11:44:35
2016-03-29 14:58:54,161 [main] INFO org.apache.pig.Main - Logging error messages to:
/home/2013/jdsilv2/MyStuff/hd/pig_1459277934159.log
2016-03-29 14:58:54,182 [main] INFO org.apache.pig.impl.util.Utils - Default bootup file
/home/2013/jdsilv2/.pigbootup not found
2016-03-29 14:58:54,737 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is
deprecated. Instead, use mapreduce.jobtracker.address
2016-03-29 14:58:54,738 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is
deprecated. Instead, use fs.defaultFS
2016-03-29 14:58:54,738 [main] INFO org.apache.pig.backend.hadoop.executionengine.HExecutionEngine -
Connecting to hadoop file system at: hdfs://cs421-hd1.cs.mcgill.ca:9000
2016-03-29 14:58:55,544 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is
deprecated. Instead, use fs.defaultFS
grunt> fs -ls
Found 1 items
-rw-r--r-- 3 jdsilv2 supergroup 4647 2016-03-29 14:54 mapredsetup.txt
grunt> fs -rm mapredsetup.txt
2016-03-29 14:59:07,975 [main] INFO org.apache.hadoop.fs.TrashPolicyDefault - Namenode trash configuration:
Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted mapredsetup.txt
grunt>
```

To copy a file from HDFS to local file system

```
grunt> copyToLocal /data2/mydata.csv /tmp/mylocalcopy.csv
grunt>
```

To see a list of files in HDFS use `ls`

```
grunt> ls
```

To see contents stored in a file in HDFS, use `cat` commands

```
grunt> cat part-r-00000
```

To see more commands, type `?` at the grunt prompt. ( There's quite a bit of mismatch between what the listing provides and what grunt actually supports, so don't get engrossed in some of the complex options, they most likely are not implemented yet.)

```
grunt> ?
```

Due to a bug in the framework, some of the website links generated by the job history / resource manager web interface, will not have fully qualified hostnames and can as a result cause your browser to not find the webpage when you click on it due to DNS failure. This can be addressed by either editing the offending link in the browser to include the full name of the host, or including the IP address mapping in the hosts file of your laptop/computer from which you are using the browser as shown below (The location of hosts file is different for different operating systems, this is **/etc/hosts** for Mac and Linux operating systems . For windows it's usually **%SystemRoot%\System32\drivers\etc\hosts**

```
132.206.51.211 winter2022-comp421-1.CS.McGill.CA winter2022-comp421-1
132.206.51.212 winter2022-comp421-2.CS.McGill.CA winter2022-comp421-2
132.206.51.213 winter2022-comp421-3.CS.McGill.CA winter2022-comp421-3
132.206.51.214 winter2022-comp421-4.CS.McGill.CA winter2022-comp421-4
132.206.51.215 winter2022-comp421-5.CS.McGill.CA winter2022-comp421-5
```

To terminate grunt shell in interactive mode, you can type `quit`;

```
grunt> quit;
```

If you are executing a pig script, you can do CTRL+C to terminate it.

**It should also be noted that in general MapReduce jobs will run a LOT longer than typical database queries.** The example pig script will take around 2 minutes and this can become longer as the number of total jobs in the system increases and your job would end up in the wait queue. There are also individual user capacity limits to ensure one user does not hog all the system resources. So if you submit multiple jobs at the same time, it might start slowing down your own throughput instead. Hence we strictly advise not to submit more than one pig script at a time. Ignoring our repeated warning can result in your id being suspended !

## How to start writing your script

We encourage you to start writing your script by typing in commands one after the other in the grunt shell, so that you get immediate feedback from pig if there is an error in your statement. However, the job history manager records the commands submitted from grunt shell as “DefaultJobName” so you may not be able to easily tell apart your scripts in the history manager UI. You can explicitly set a job name for the set of commands you submit as shown below.

```
grunt> set job.name 'YourSOCSID_JobXX';
```

Once you have all the commands working as desired, you can write them into a single script file and execute them together. **It should also be noticed that pig submits the jobs to Hadoop ONLY when it encounters a STORE or DUMP command.**

## Support and Questions

**If you have questions regarding the setup, please post it in Piazza and tag assignments/a3. Do not email the cs helpdesk with issues you have on the MapReduce cluster, they are not responsible for the cluster setup.**

## Useful links

Basics

<https://pig.apache.org/docs/r0.15.0/basic.html>

How to generate some very useful diagnostic / informational outputs that you can leverage in writing answers to the general questions asked.

<https://pig.apache.org/docs/r0.15.0/test.html>

Oreilly’s Programming Pig ebook (Because you are insanely obsessed with Pig)

<https://mcgill.on.worldcat.org/oclc/992150957>